

Installing, Starting, and Using H2O Steam

Copyright H2O, Inc. All rights reserved.

May 2016: Alpha Edition

Installing, Starting, and Using H2O Steam

Published by H2O.ai, Inc.

2307 Leghorn St.

Mountain View, CA 94043

©2016 H2O.ai, Inc. All Rights Reserved.

May 2016: Alpha Edition

Photos by ©H2O.ai, Inc. All copyrights belong to their respective owners. While every precaution has been taken in the preparation of this book, the publisher and authors assume no responsibility for errors or omissions, or for damages resulting from the use of the information contained herein. Printed in the United States of America.

Contents

Installing, Starting, and H2O Using Steam	4
Requirements	4
Installation	4
Starting Steam	5
Using Steam	5
Adding an Engine	5
Starting a Cloud	6
Adding a Model	9
Viewing Models	9
Deploying Models	10
Using Steam with H2O Flow	10
Stopping Steam	11

Installing, Starting, and H2O Using Steam

This document is provided for external users and describes how to install, start, and use Steam on a YARN cluster.

Note: This installation should only be performed on a YARN edge node.

Requirements

- Web browser with an Internet connection
- Steam tar for OS X or Linux
- JDK 1.7 or greater
- H2O AutoML package (for example, **automl-hdp2.2.jar**)

Installation

Perform the following steps to install Steam.

1. Open a terminal window and SSH into your YARN edge node.
2. Retrieve the Steam release for your OS.

Linux

```
s3cmd get s3://steam-release/steamY-master-linux-amd64.tar.gz
```

OS X

```
s3cmd get s3://steam-release/steamY-master-darwin-amd64.tar.gz
```

Note: This saves the Steam package (tar file) into your current folder.

3. Untar the Steam package.

Linux

```
tar -xzf steam--linux-amd64.tar.gz
```

OS X

```
tar -xsf steam--darwin-amd64.tar.gz
```

Starting Steam

Now that Steam is installed, perform the following steps to start and use Steam. Note that two terminal windows will remain open: one for the Jetty server and one for Steam.

1. Change directories to your Steam directory, then set up the Jetty server using one of the following methods:

```
cd steam--linux-amd64
java -jar var/master/assets/jetty-runner.jar --port 8811 var/master/assets/ROOT.war
```

2. Open another terminal window, then start Steam on the master server (the compilation server that will run the POJOs). For example, the following commands will start Steam on localhost. Note that the port value must match the port specified when starting the Jetty server.

```
./steam serve master --compilation-service-address="localhost:8811"
```

Note: You can view all available options for starting Steam using `./steam help serve master`

You will see a message similar to the following when Steam starts successfully.

```
2016/04/28 13:34:56 steam v build 2016-04-28T20:15:00+0000
2016/04/28 13:34:56 Working directory: /home/seb/steam--linux-amd64/var/master
2016/04/28 13:34:56 WWW root: /home/seb/steam--linux-amd64/var/master/www
2016/04/28 13:34:57 Priming datastore for first time use...
2016/04/28 13:34:57 Datastore location: /home/seb/steam--linux-amd64/var/master/db
/steam.db
2016/04/28 13:34:57 Web server listening at 172.16.2.182:9000
2016/04/28 13:34:57 Point your web browser to http://172.16.2.182:9000/
```

Using Steam

In a Web browser, navigate to the Steam Web server (for example, <http://172.16.2.182:9090>).

Adding an Engine

An empty Steam UI will display. Before performing any tasks, you must first add an Asset. In this case, the **automi-hdp2.2.jar** file (or similar) provides the engine necessary for Steam to run AutoML jobs.

1. Click the **Assets** icon on the left navigation panel, then select **Add Engine**.
2. Browse to the **automi-hdp2.2.jar** file on your local machine, then click **Upload**.

Note: This file was added during the “Install Steam” steps.

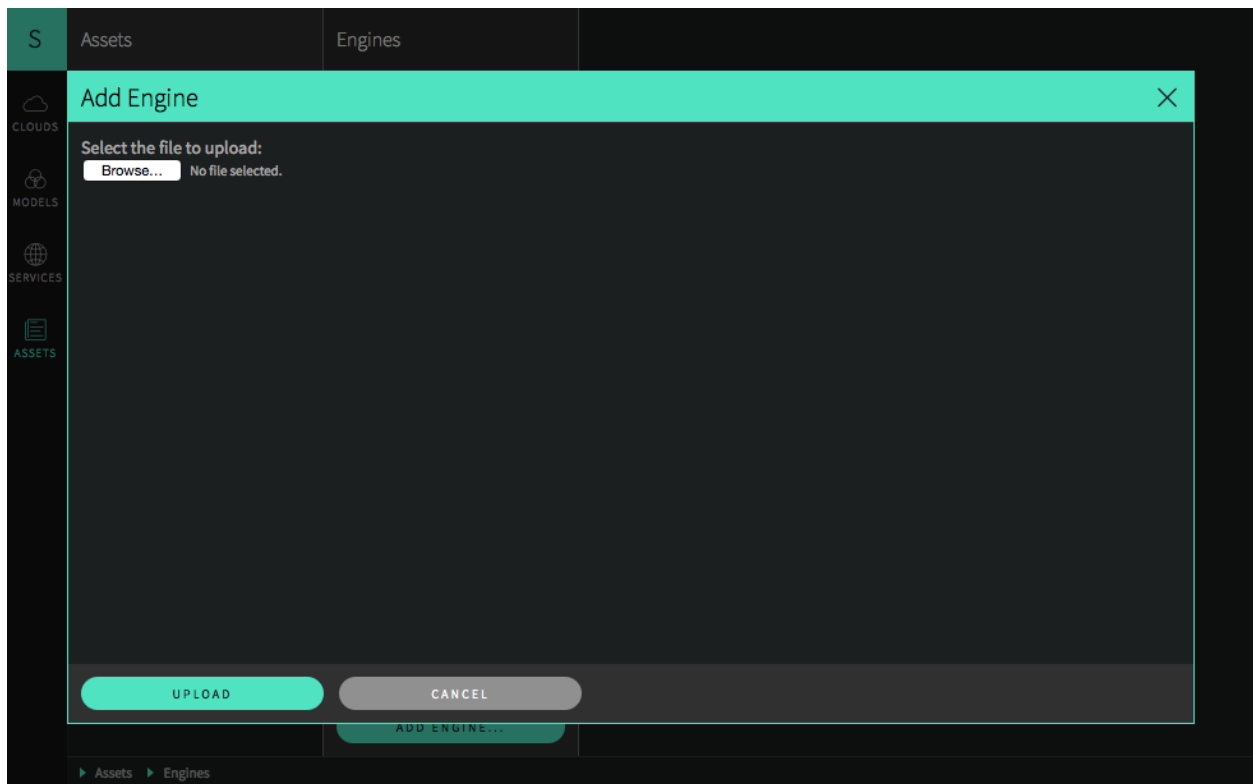


Figure 1: Adding an Engine

Starting a Cloud

Clouds can be configured after the engine asset was successfully added.

1. Click the **Clouds** icon on the left navigation panel, then select **Start a Cloud**.
2. Enter/specify the following information to set up your cloud:
 - a. A name for the cloud
 - b. The version of H2O that will run on the cloud.
 - c. The number of nodes on the cloud.
 - d. The amount of memory available on each node. Be sure to include the unit ("m" or "g").
3. Click **Start Cloud** when you are finished.

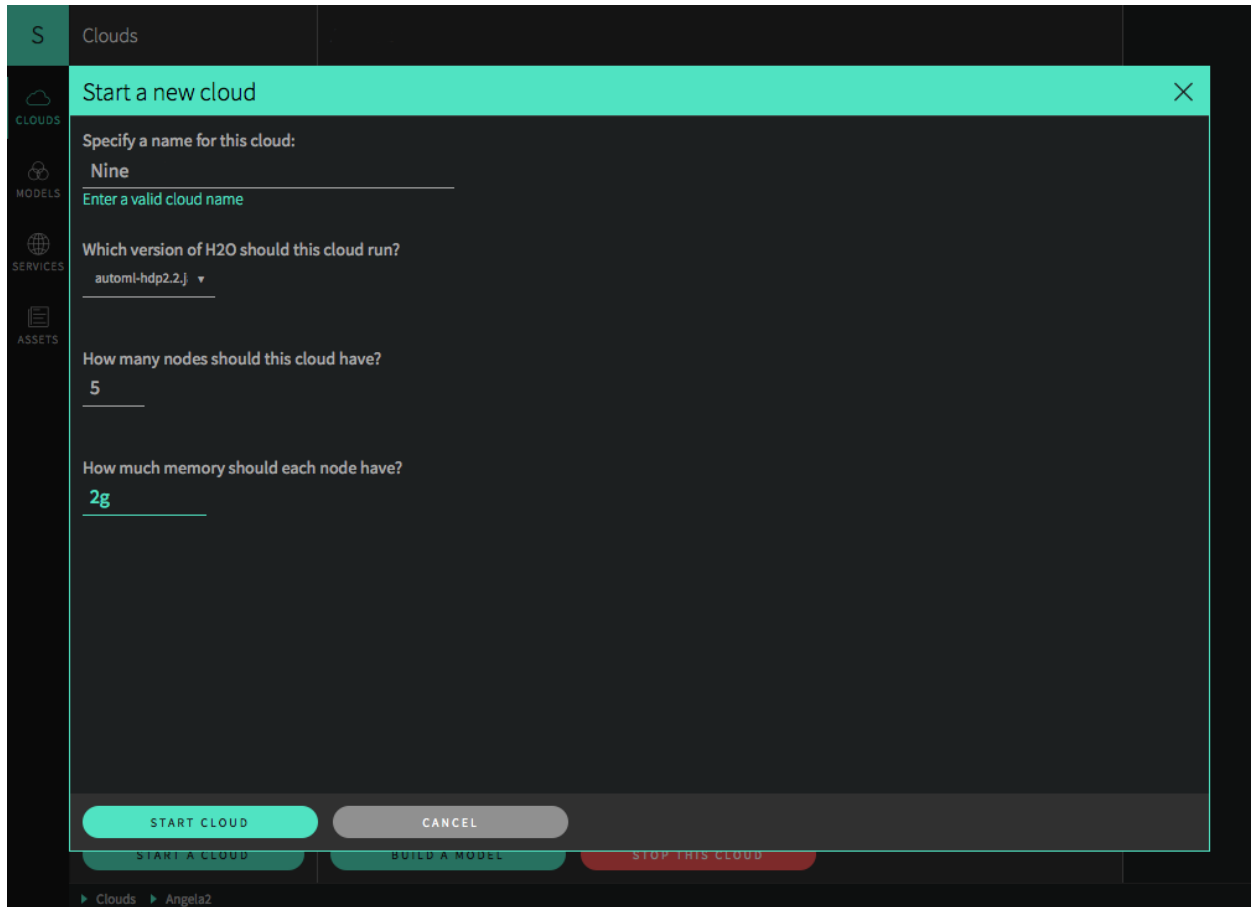


Figure 2: Adding a cloud

The Cloud Details page opens upon successful completion. This page shows the cloud configuration information and includes a link to the H2O Flow URL. From this page, you can begin building your model.

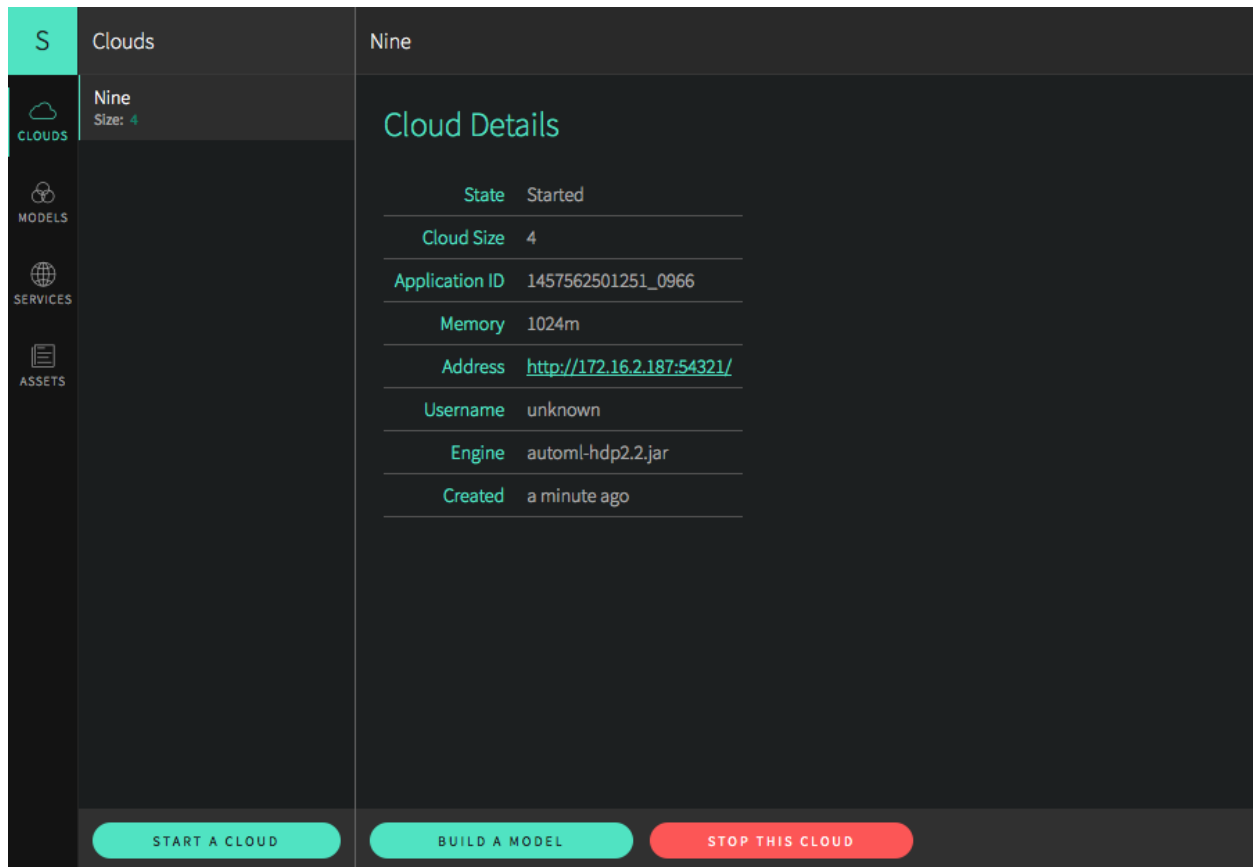


Figure 3: Cloud details

Note: You can view a stream of the cloud creation log in the terminal window that is running Steam. In the UI, Steam will respond with an error if the cloud configuration is incorrect (for example, if you specify more nodes than available on the cluster).

Note: On the Cloud Details page, the **Memory** field shows the memory for each node rather than the total cloud memory.

Adding a Model

Models are created from the Cloud Details page. When building a model, you will need to provide the location of the dataset that you will use as well as the response column.

1. Click the **Build a Model** button on the bottom of the Cloud Details page.
2. Enter a path for the dataset that you want to use to build the model.
Note: If you choose to use a local dataset, then that dataset must reside in the same folder/path on each node in the cluster.
3. Specify the column that will be used as the response column in the model.
4. Specify the maximum run time in seconds. H2O will return an error if the model build stalls after this threshold is reached.
5. Click **Start Building** when you are finished.

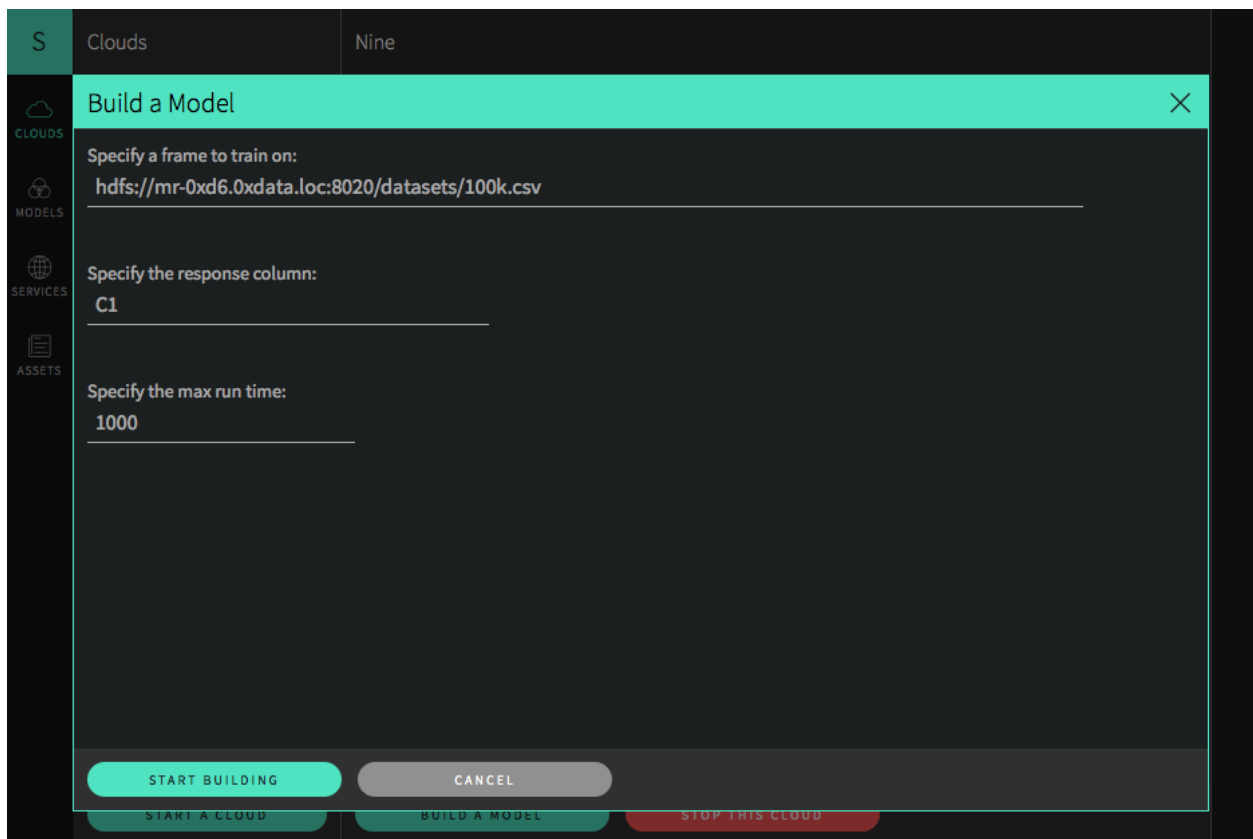


Figure 4: Building a model

Viewing Models

Click the **Models** icon on the left navigation panel to view models that were successfully created.

These models are processed using H2O's AutoML algorithm, which determines the best method to use to build the model. The model name includes this method. So, for example, if Steam

returns a model named “DRF_model. . .”, then this indicates that DRF was the algorithm that provided the best result.

Deploying Models

After a model is built, the next step is to deploy the model in order to make/view predictions.

1. Click the **Models** icon on the left navigation panel.
2. Select the model that you want to use, then click the **Deploy this Model** button on the bottom of the page.
3. Specify the port to use for the scoring service.

Note: Steam will return an error if you specify a port that is already being used.

4. Click **Deploy** when you are done.

Making Predictions

Successfully deployed models can be viewed on the Services page. On this page, click the Endpoint link to begin making predictions.

Using Steam with H2O Flow

As with other H2O products, Flow can be used alongside Steam when performing machine learning tasks.

On the Cloud Details page, click the Address link to open H2O Flow in a new tab.

Note: Refer to the H2O Flow documentation for information on how to use Flow.

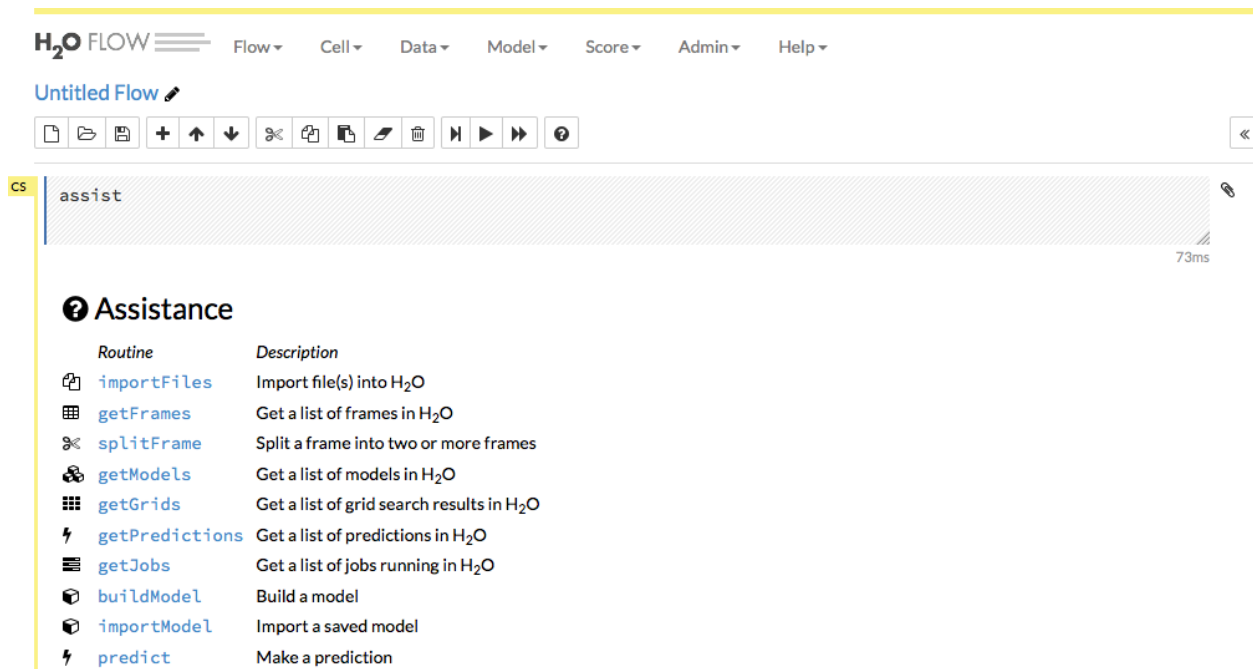


Figure 5: H2O Flow UI

Stopping Steam

When you are finished, use the following process to safely shut down Steam:

1. Stop all running services.
2. Stop all running clouds.